

# Assignment 4: Regression

## A. Simple Linear Regression

*a) Import the dataset in Python environment.*

In [1]:

```
import numpy as np
import pandas as pd
import seaborn as sb
import matplotlib.pyplot as plt
from sklearn import preprocessing, svm
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
```

In [2]:

```
df = pd.read_csv('USA_Housing.csv')
```

*b) Make a new dataset considering only 'Avg. Area House Age'; and 'Price' attributes.*

In [3]:

```
df_binary = df[['Avg. Area House Age', 'Price']].copy()
df_binary.columns = ['Avg. Area House Age', 'Price']
df_binary.head()
```

Out[3]:

	Avg. Area House Age	Price
0	5.682861	1.059034e+06
1	6.002900	1.505891e+06
2	5.865890	1.058988e+06
3	7.188236	1.260617e+06
4	5.040555	6.309435e+05

*c) Set a variable X equal to the 'Avg. Area House Age' feature of the given dataset and a variable y equal to the "Price" column.*

In [4]:

```
x=df['Avg. Area House Age']  
y=df['Price']  
x
```

Out[4]:

```
0      5.682861  
1      6.002900  
2      5.865890  
3      7.188236  
4      5.040555
```

...

```
4995    7.830362  
4996    6.999135  
4997    7.250591  
4998    5.534388  
4999    5.992305
```

Name: Avg. Area House Age, Length: 5000, dtype: float64

In [5]:

```
y
```

Out[5]:

```
0      1.059034e+06  
1      1.505891e+06  
2      1.058988e+06  
3      1.260617e+06  
4      6.309435e+05
```

...

```
4995    1.060194e+06  
4996    1.482618e+06  
4997    1.030730e+06  
4998    1.198657e+06  
4999    1.298950e+06
```

Name: Price, Length: 5000, dtype: float64

**d) Split the new dataset into the Training set and Test set such that Test set consists 1/3 of total records using `model_selection.train_test_split` from `sklearn`.**

In [6]:

```
X = np.array(df_binary['Avg. Area House Age']).reshape(-1, 1)  
Y = np.array(df_binary['Price']).reshape(-1, 1)  
df_binary.dropna(inplace = True)  
x_train,x_test,y_train,y_test=train_test_split(X,Y,test_size=0.33)
```

**e) Train the Simple Linear Regression model based on the Training set after importing `LinearRegression` from `sklearn.linear_model`.**

In [7]:

```
lr = LinearRegression()  
lr.fit(x_train,y_train)
```

Out[7]:

```
LinearRegression()
```

**f) Display the model's coefficients after train the model and hence write the final regression model.**

In [8]:

```
print(lr.score(x_test,y_test))
```

```
0.19740199111771872
```

In [9]:

```
lr.coef_
```

Out[9]:

```
array([[161628.45007805]])
```

In [10]:

```
lr.intercept_
```

Out[10]:

```
array([264376.53968167])
```

In [11]:

```
m=lr.intercept_  
b=lr.coef_
```

In [12]:

```
print("Y=",m,"*X+",b)
```

```
Y= [264376.53968167] *X+ [[161628.45007805]]
```

**g) Predict the house price of the Test set data and display them.**

In [13]:

```
pred=lr.predict(x_test)
```

In [14]:

```
print(pred)
```

```
[[1306425.03128295]  
 [1256462.13592285]  
 [1168569.88135383]  
 ...  
 [1175667.64798037]  
 [1122970.9181027 ]  
 [1085855.73974469]]
```


***h) Visualize the Training set results using scatter plot along with regression line.***

In [15]:

```
plt.scatter(x_train,y_train)  
plt.plot(x_test,pred,color='red')
```

Out[15]:

```
[<matplotlib.lines.Line2D at 0x1b32b228670>]
```




***i) Visualize the Test set results using scatter plot along with regression line.***

In [16]:

```
plt.scatter(x_test,y_test)  
plt.plot(x_test,pred,color='red')
```

Out[16]:

[<matplotlib.lines.Line2D at 0x1b32d49c700>]



## B. Multiple Linear Regression

**a) Use the dataset which you have already imported in Python environment in Question a of Section A above.**

In [17]:

```
df1=pd.read_csv("USA_Housing.csv")
```

**b) Check the head of your Dataset, and also check out its info() and describe() methods over the dataset.**

In [18]:

```
df1.head()
```

Out[18]:

	Avg. Area Income	Avg. Area House Age	Avg. Area Number of Rooms	Avg. Area Number of Bedrooms	Area Population	Price	Address
0	79545.45857	5.682861	7.009188	4.09	23086.80050	1.059034e+06	208 Michael Ferry Ap 674\nLaurabury, N 3701
1	79248.64245	6.002900	6.730821	3.09	40173.07217	1.505891e+06	188 Johnson View Suite 079\nLak Kathleen, CA
2	61287.06718	5.865890	8.512727	5.13	36882.15940	1.058988e+06	9127 Elizabe Stravenue\nDanieltow WI 06482
3	63345.24005	7.188236	5.586729	3.26	34310.24283	1.260617e+06	USS Barnett\nFPO A 4482
4	59982.19723	5.040555	7.839388	4.23	26354.10947	6.309435e+05	USNS Raymond\nFP AE 0938

In [19]:

```
df1.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 5000 entries, 0 to 4999  
Data columns (total 7 columns):  
#   Column                                Non-Null Count  Dtype  
---  -  
0   Avg. Area Income                      5000 non-null   float64  
1   Avg. Area House Age                   5000 non-null   float64  
2   Avg. Area Number of Rooms             5000 non-null   float64  
3   Avg. Area Number of Bedrooms          5000 non-null   float64  
4   Area Population                       5000 non-null   float64  
5   Price                                 5000 non-null   float64  
6   Address                               5000 non-null   object  
dtypes: float64(6), object(1)  
memory usage: 273.6+ KB
```

In [20]:

```
df1.describe()
```

Out[20]:

	Avg. Area Income	Avg. Area House Age	Avg. Area Number of Rooms	Avg. Area Number of Bedrooms	Area Population	Price
<b>count</b>	5000.000000	5000.000000	5000.000000	5000.000000	5000.000000	5.000000e+03
<b>mean</b>	68583.108984	5.977222	6.987792	3.981330	36163.516039	1.232073e+06
<b>std</b>	10657.991214	0.991456	1.005833	1.234137	9925.650114	3.531176e+05
<b>min</b>	17796.631190	2.644304	3.236194	2.000000	172.610686	1.593866e+04
<b>25%</b>	61480.562390	5.322283	6.299250	3.140000	29403.928700	9.975771e+05
<b>50%</b>	68804.286405	5.970429	7.002902	4.050000	36199.406690	1.232669e+06
<b>75%</b>	75783.338665	6.650808	7.665871	4.490000	42861.290770	1.471210e+06
<b>max</b>	107701.748400	9.519088	10.759588	6.500000	69621.713380	2.469066e+06

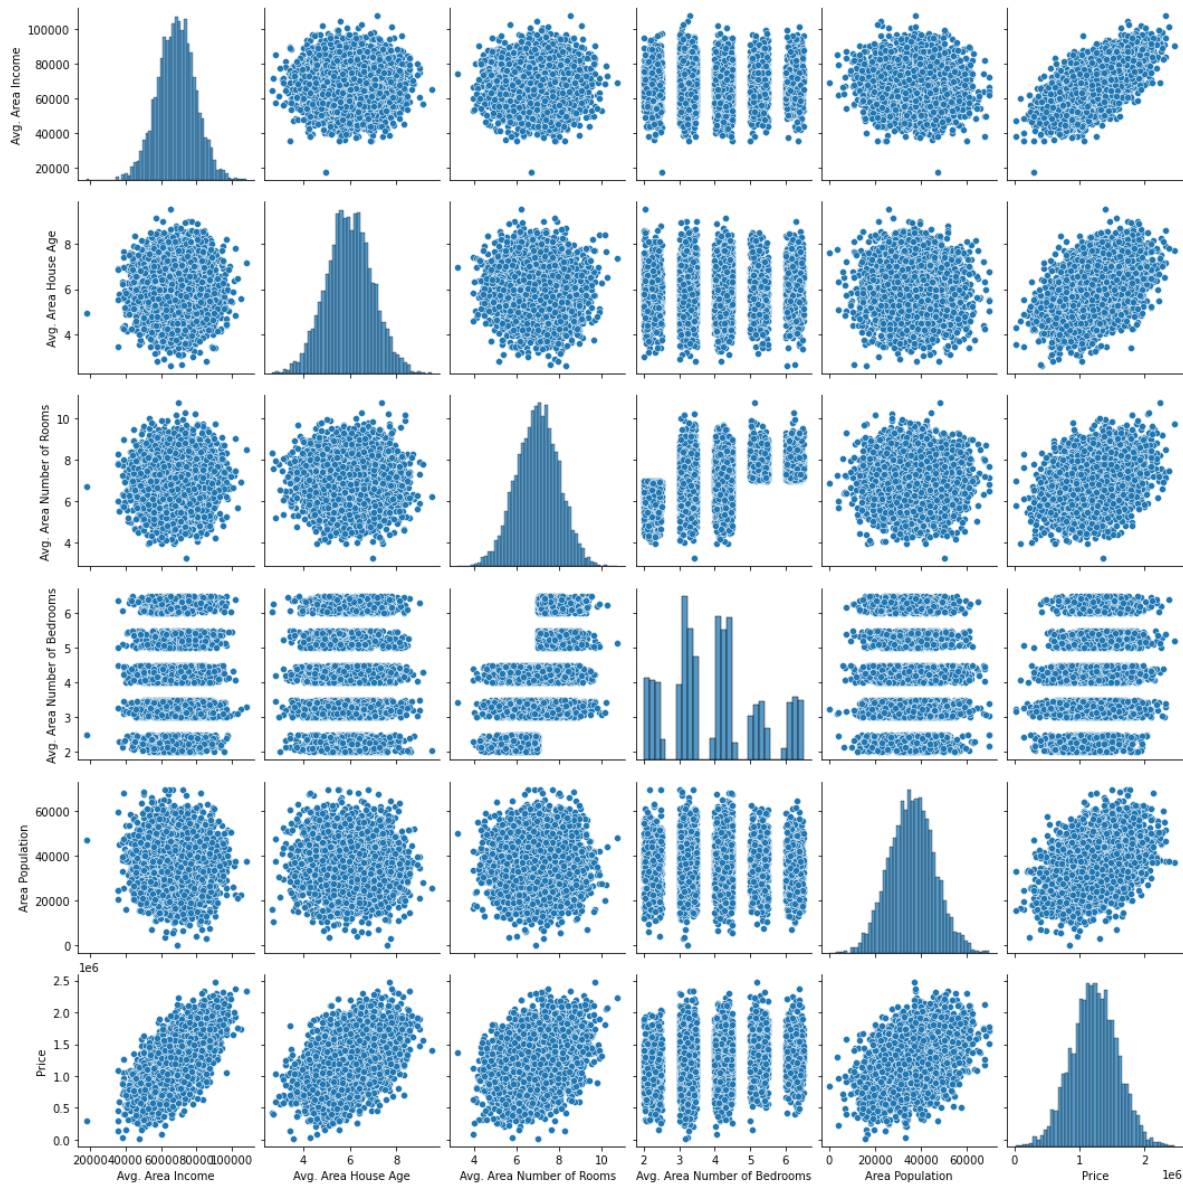
**c) Explore the types of relationships across the entire data set using 'pairplot' method of seaborn and comment on that.**

In [21]:

```
sb.pairplot(df1)
```

Out[21]:

&lt;seaborn.axisgrid.PairGrid at 0x1b32d71d1f0&gt;





**d) Set a variable  $X$  equal to the numerical features of the given dataset and a variable  $y$  equal to the "Price" column.**

In [22]:

```
X = df1[['Avg. Area Income', 'Avg. Area House Age', 'Avg. Area Number of Rooms',  
        'Avg. Area Number of Bedrooms', 'Area Population']]  
Y=df1['Price']
```

**e) Split the data into training and testing sets using `model_selection.train_test_split` from `sklearn` such that Test set consists 30% of total data.**

In [23]:

```
x_train,x_test,y_train,y_test=train_test_split(X,Y,test_size=0.30)
```

**f) Train the Linear Regression model based on the Training data after importing `LinearRegression` from `sklearn.linear_model`.**

In [24]:

```
lm1=LinearRegression()  
lm1.fit(x_train,y_train)
```

Out[24]:

```
LinearRegression()
```

**g) Print the coefficients of the trained model. How can you interpret these coefficients?**

In [25]:

```
print(lm1.score(x_test,y_test))
```

```
0.9204289208842987
```

In [26]:

```
print(lm1.intercept_)
```

```
-2631251.8071228913
```

In [27]:

```
coeff_df = pd.DataFrame(lm1.coef_, X.columns, columns=['Coefficient'])
coeff_df
```

Out[27]:

	Coefficient
Avg. Area Income	21.588120
Avg. Area House Age	165100.159433
Avg. Area Number of Rooms	119746.635081
Avg. Area Number of Bedrooms	2765.916748
Area Population	15.146601

**h) Predict the house price of the Test set data and display them.**

In [28]:

```
prediction=lm1.predict(x_test)
print(prediction)
```

```
[1225074.00144784 1349362.37857084 991051.61288649 ... 1098298.67538327
 1187409.46602538 1363950.36658695]
```


**i) Create a scatterplot of the real test values versus the predicted values.**

In [29]:

```
plt.scatter(y_test,prediction,edgecolor='red')
```

Out[29]:

```
<matplotlib.collections.PathCollection at 0x1b32fde31f0>
```



**j) Calculate the Mean Absolute Error, Mean Squared Error, and the Root Mean Squared Error to evaluate our model performance after importing metrics from sklearn.**

In [30]:

```
from sklearn import metrics
print('MAE:', metrics.mean_absolute_error(y_test, prediction))
print('MSE:', metrics.mean_squared_error(y_test, prediction))
print('RMSE:', np.sqrt(metrics.mean_squared_error(y_test, prediction)))
```

MAE: 81491.64716425468  
 MSE: 10003687394.799738  
 RMSE: 100018.43527470193

**k) Plot a histogram of the residuals. [Use either seaborn distplot, or just plt.hist()].**

In [31]:


```
sb.distplot((y_test - prediction), bins = 50, hist_kws=dict(edgecolor="red", linewidth=1))
```

C:\Users\hp\anaconda3\lib\site-packages\seaborn\distributions.py:2551: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

warnings.warn(msg, FutureWarning)

Out[31]:

<AxesSubplot:xlabel='Price', ylabel='Density'>



In [ ]: